

Hardware

CPU's

GPU's / accelerators

Network

Disk

Memory RAM / caches

Software

algorithms

PL / runtime sys.

OS

architectures

MT event-based
MP

Making force of Perf in Data Analytics Framework ©

blocked time analysis

infinite resources

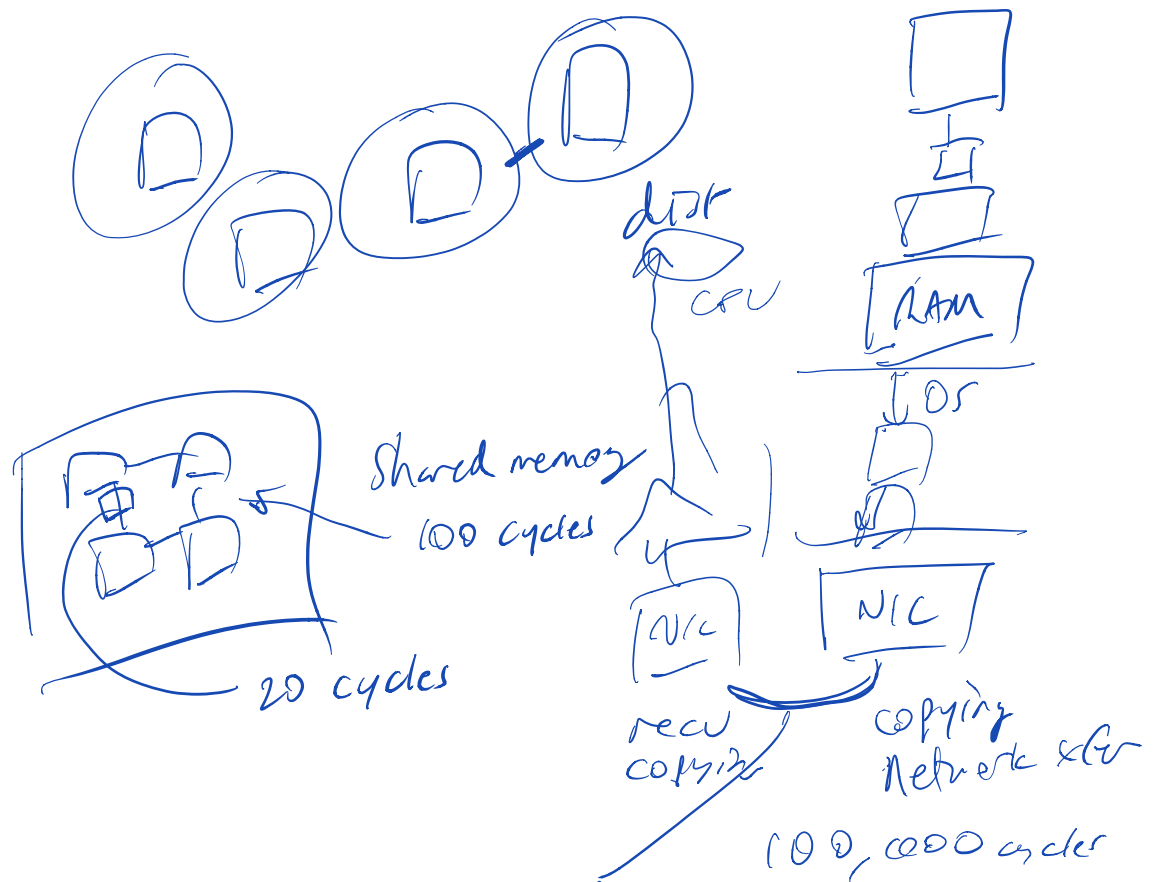
Spark SQL

CPU "bound"


I/O bound

disk bound


network bound



~~RDMA~~
remote direct memory access



10,000 cycles




141

Spark vs Hadoop

not always better

① data too large to fit in RAM
workload 2TB
RAM capacity 100 GB
"paging"


② No situation





perf. workload dependent

GC remains a problem

Google → "MapReduce" C++
~~GC~~ GC → safety




GC pause times
space overhead




~ 3-5x more than
using malloc/free

full heap GC
→ pause times

Chrome v8 GC
 PLDT 16
 "Metronome"
 Bacon



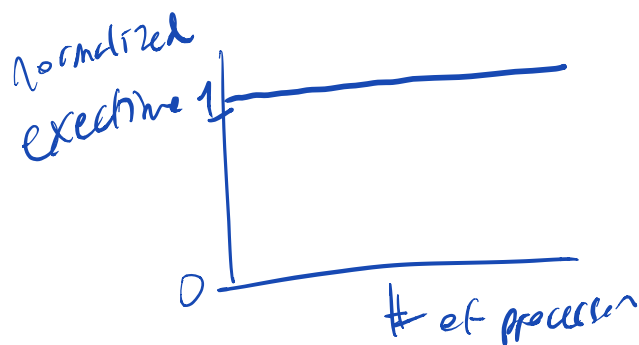


Project Tungsten — manual
 memory mgmt in Spark!
 goal → set nil of GC



$$\sum_{i=1}^{1e^9} i$$

```

s=0
for(i=1; i<1e9; i++) {
  s+=i
}
return s;
  
```



$$S = \frac{1e9 \times (1e9 + 1)}{2}$$

baseline: fastest state-of-the-art version
on one node



COST Config. Outperforms a Single Thread

$\max(\infty, \# \text{ of cores required to match perf. sit. version})$

problems -

too big to fit in L1m
too big to fit on disk

~ Gustafson's Law
algorithms for single thread
* also. for dist. comp.
fault tolerance (solvable)



Domain-Specific Languages

SQL

Excel

"analytics"

libraries \approx DSL



Java
JS


Python

general-purpose

s.t. \rightarrow s.t.! ONE THREAD

multiple processors





Space-filling Curve

Z-curve

