

NoSQL
SQL

MapReduce
databases

triples
(key, value)

"triples"

OO databases \approx NoSQL "objects"

XML



"trees"

graphs



adjacency matrix



Symmetric
Circulant
diag

undirected
directed

weighted




vertex-edge notation

directed
A \rightarrow B
B \rightarrow A

undirected — symmetric
space,
FB social networks
internet

directed — asymmetric
internet
Twitter



/weighted

static — always unchanging (cities)

dynamic — potentially changing



vertex vertex
 (1, 2)
 (1, 3)
 (2, 3)
 (2, 4)

implicit graph representation -
 (edges)
 edge-centred



(traffic network)
 social network
 deleted edges
 new edges
 new vertices (users)
 deleted vertices

- edge list
 (2, vertices)
- ① [2, 3]
 - ② [3]
 - ③ [1, 2]
 - ④ [2]



1 2 3 4 5 6 7

$$\begin{array}{c}
 1 \\
 2 \\
 3 \\
 4 \\
 5 \\
 6 \\
 7
 \end{array}
 \left|
 \begin{array}{ccccccc}
 1 & 1 & 0 & 0 & 0 & 0 & 0 \\
 & 1 & 1 & 0 & 0 & 0 & 0 \\
 & & 1 & 1 & 0 & 0 & 0 \\
 & & & 1 & 1 & 0 & 0 \\
 & & & & 1 & 1 & 0 \\
 & & & & & 1 & 1 \\
 & & & & & & 1
 \end{array}
 \right|$$

Sparse
lot of zeros

dense matrix $(N \times N) \rightarrow N^2$

sparse matrix $(N \times N)$

1 entry $O(1)$

Coord + value


$(1, 1,)$ $(1, 1, 4)$
 $(1, 2,)$ —
 $(1, 3,)$ —
 $(1, 4,)$ —

$$\left|
 \begin{array}{cccccc}
 1 & 0 & 1 & 0 & 1 & 0 \\
 1 & 0 & 1 & 0 & 1 & 0 \\
 & & & & & \\
 & & & & & \\
 & & & & &
 \end{array}
 \right|
 \begin{array}{l}
 (1, 1, 1) \\
 (1, 3, 1) \\
 (1, 5, 1) \\
 \vdots
 \end{array}$$

Social net:

N vertices

$\ll N^2$ edges



5 vertices

$$\frac{N(N-1)}{2} = 10$$
$$O(N^2)$$

Complete graph



sparse representation

- run-length encoding

sparse matrix

(suboptimal for graphs)

- (vertex, vertex) $O(E)$ $\sim 2E$

- edge lists $O(V+E)$ $\sim V+E$

V vertices

E edges

if $V \ll E$
 $\approx E$

edge list $\begin{matrix} \textcircled{1} \\ \textcircled{2} \\ \vdots \end{matrix} \begin{matrix} [edges] \\ [edges] \\ \vdots \end{matrix}$ ^{weights} key, [values]

$(v, (e^w)) \Rightarrow (1, [2:3, 3:10, 4:100])$
_{weight(e)}

\Downarrow

$\begin{pmatrix} \textcircled{2}, 2, 3 \\ \textcircled{3}, 10 \\ \textcircled{4}, 100 \end{pmatrix}$

— mismatch
between
problem

2 interface
abstraction

$(2, \dots)$

— low utilization


— performance

iteratively
executing
mapreduce

\Rightarrow BAD
 (I/O)

— locality — insensitive

framework
oblivious
to fact these
are graphs



Friend
of
a friend

graph databases

Input

vertices edge lists

"Combinatorial
academics"

write only

LPU

Salami
string

locality

which vertex so where

vertex-centred

"Think like a vertex"

all edges are with each vertex

problems: Justin Broder

Ellen DeGere

selfie

avoid messager

=> local

(fault tolerance)

too many!


IN RAM

BSP bulk-synchronous parallelism


Valiant 1988



programming model
perf guarantees / cost model

~~PRAM model~~ academic



MPI Message Passing Interface





GENERAL
Uncoordinated
parallel
execution
MT, MP

broadcast(all)
group communication (some)
point-to-point

mpi_receive(2, asynch
Synch
(MPI_ANY

problems
deadlock

SPMD

single program
multiple data

```
if (mpi_self == 0) {  
    mpi_recv(2, b)
```

No shared memory
(distributed)

```
if (mpi_self == 0)  
else {  
    mpi_recv(2, b)  
}
```

```
if (mpi_self() == 0) {  
    // master  
} else  
    // worker  
}
```



```
if (id == 0)
    mpi_recv(Any)
    print result
```

```
else
    mpi_send(self, 0)
```

"race" (message passing)
asynch ring



determinism ☺

barriers ☹

synchronous





as long
as no deadlock

BSP — ^{comm} no cycles \Rightarrow no deadlocks



1 ms

1



Central
limit
theorem

Sum of ~~add~~ independent identical distributed variables

\Rightarrow approx. normally distributed

Random \rightarrow load balance