


multiprocessing - difficult
- doesn't solve the problem

Demand scaling done!



parallel programming

shared memory (threads)
distributed memory
- (processes)



$x = 12$
 $z = y$

explicit
communication
message-passing

"duals"
of
each
other


`machine(p2).sendMessage(p2, "get", y)`

(Lauer & Needham, 1979)

threads

processes

$x = 12;$
 $t = \text{new Thread}();$
 $\text{System.out.println}(x);$
 $t.\text{wait}();$



$x = 13;$ - local variables ("stack")
unshared
- thread local value

C++14
pthread_create
pthread_join

race condition
↳ nondeterminism

deterministic function: $f(x) = f(x)$

non-deterministic function: $f() \Rightarrow \{12, 13\}$

Control
ordering


- wait for threads
- synchronization of
 - prevents concurrency
 - orders events

synchronized
void foo()

$\Rightarrow \text{this.lock}();$

$\left(\begin{array}{l} \text{foo}() \\ \text{this.unlock}(); \end{array} \right)$

lock =
mutual
exclusion




C++

```
this->lock();
this->lock();
```

deadlock
Edsger Dijkstra


this.lock()
this.unlock() } synchronized

wait
notify



```
class Foo {
    public int x;
    void f1() { x=12; }
    void f2() { x=13; }
}
```

```
synchronized(this) {
    x=12;
}
```



thread 1
f.f1();

thread 2
f.f2();


x = 12;
z = 13;
z = 100; *happens-before*
new t1
new t2

MultiThreaded challenges:

- ① correctness
- ② performance

Heisenbug


non-determinism
potential interleavings





Sequential

s1: main() \downarrow
 s2: x = (2)
 s3: y = (3)
 s4: z = 1000000
 s5:



s1: main
 s2: new Thread f;
 s3: new Thread g;
 s4: waitforall();

$$T_1 = 9 \text{ units "statements"}$$

$$\max \{ 3, 4, 7 \} = 7$$


Critical path

longest seq. of stmts
from start to finish

$$T_{\infty} = 7$$


critical path "span"

$$T_p \leq \underbrace{\frac{T_1}{P}} + T_{\infty}$$





$$\frac{4,000,000}{1,000,000} + 4$$

$$4 + 4 = P$$




load imbalance

serialization



server

client




wants for client request

serve it


1 thread
per client

request-
response

N clients $\rightarrow N$ threads



data



extract
word
from
html

Gustafson's
Law
in
action

inverted
index { amazon: 500, 1435, 999
android: 500, 217
prime: 500, 144

"big
data"

>> 1TB

<u>Thread₁</u>	<u>Thread₂</u>	Share?
local variables (stack variables)		X
heap?		✓
file handles		✓
globals ?		✓
code		✓

<u>Process₁</u>	<u>Process₂</u>	Share?
local variables (stack variables)		X
heap?		X
file handles		X
globals ?		X
code		✓

spawn(p)

fork()

clone()

int x = 0;

```
int r;  
if (r = fork()) {  
    x = 12;  
    print x;  
}  
else {  
    print x;  
}
```

