


Parallel DBs

multiple processors

100 computers ~ 100x faster

multiple disks




~~1 DB file → machine~~


not much I/O

may not fit

partitioning



horizontal partitioning



load balance

database(s) spread evenly


goal: hash fn "looks random"

hash


$$P[h(x) = i] \approx \frac{1}{N}$$

1...N

N machines



$h(\text{"Joe Franklin"}) \rightarrow 939715 \% \log N$
 ↗
 machine 15



(1) (1) (1) (1) . . . (1) (1) (2)

trade-off
 maximize local comp.

maximize utilization
 maximize RAM utilization
 minimize communication

DBMS Schema

NAME STRING (30)
 SSN STRING (9)
 SALARY num.

BLOBS

binary base object
 IMAGE
 variable-sized

MapReduce

Schemaless


web pages

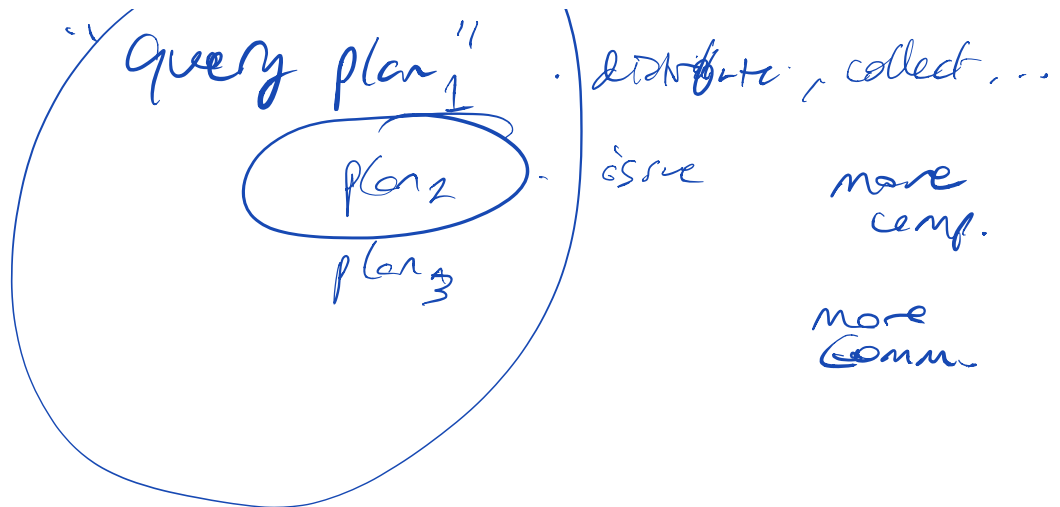
- "structured"
 hierarchical
 - "XML" / HTML



URL word
 the
 news

hash partitioned





row-based

Column-based

name	salary	age	Gender	col1	BLOB portrait

Global

name	name	name		
Salary	-	-	-	-

(column)
in
append
OK

I/O util.

$$\frac{\text{I/O I use}}{\text{I/O I do}}$$

$$\frac{16 \text{ bytes}}{4096 \text{ bytes}}$$

100%
util.

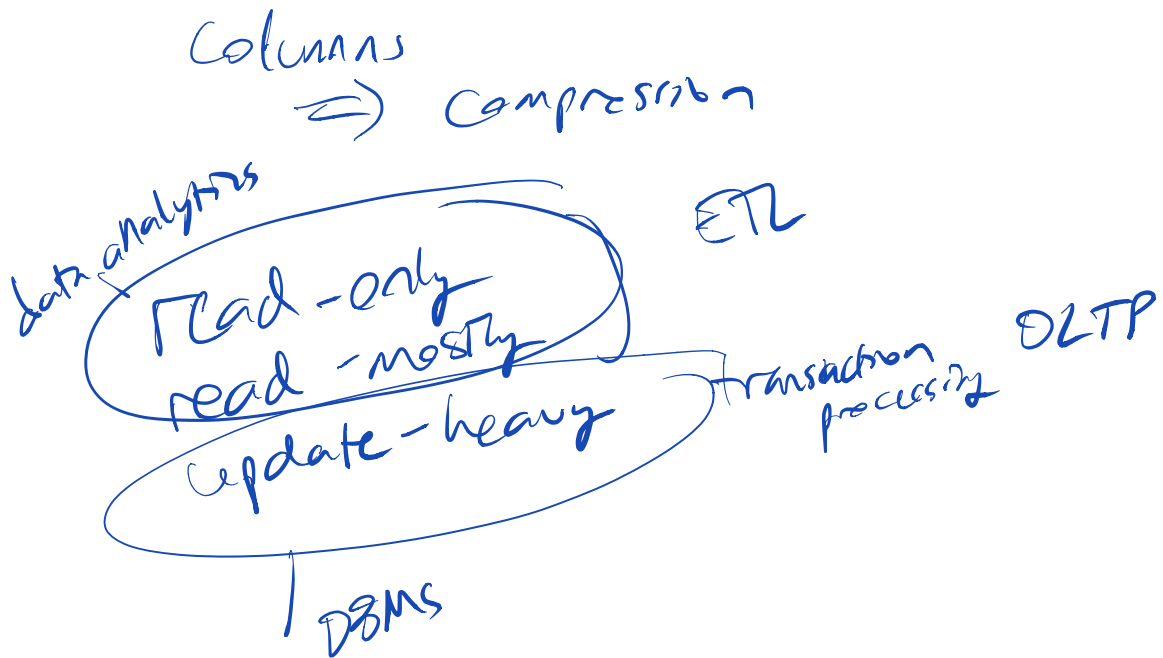
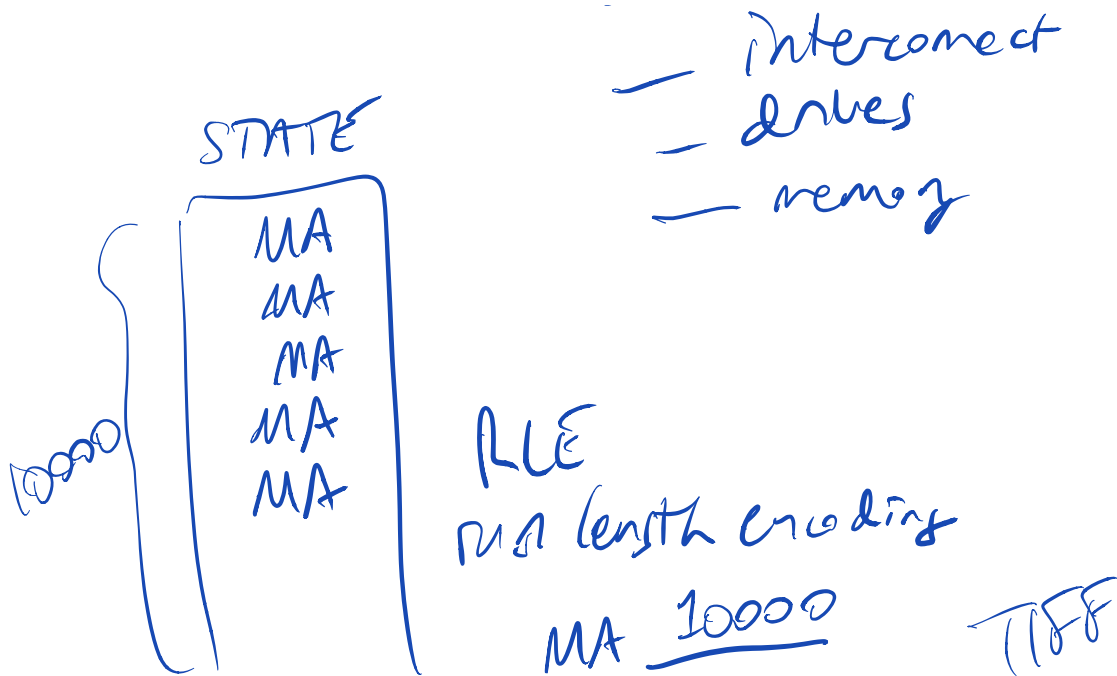
$$\frac{4096 \text{ bytes}}{4096 \text{ bytes}}$$


COL-BASED
100 column



"it depends"

- workload
- machines





UPDATE WHERE NAME = "John Franklin" — YOUR NAME GOES HERE

SALARY = SALARY + 1000000

UPDATE ~~WIDE~~ *

SALARY = SALARY * 1.02

transactions

ACID

Atomic

all or nothing

Consistent

good state \rightarrow good state


Isolated

transactions
can't see updates

Durable

saved to
persistent store

F3



ACID overkill

eventual consistency

recovery, on disk

- Locking

NoSQL — ~~did not~~ hash tables

put(k, v) ~~stores~~ stores k, v pair

$$\text{get}(k) \rightarrow v$$